

**Aim:**

Download IRIS dataset from UCI Repository The data set contains 3 classes of 50 instances each, where each class refers to a type of iris plant. One class is linearly separable from the other 2; the latter are NOT linearly separable from each other.

1. Apply Data pre-processing (Label Encoding, Data Transformation....) techniques if necessary.
2. Perform data-preparation ( Train-Test Split)
3. Apply Logistic Regression Algorithm
4. Evaluate Model.

**Importing necessary libraries**

In [46]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

**Reading the dataset**

In [24]:

```
A = pd.read_csv(r"C:\Users\RASIKA\Downloads\archive (8)\Iris.csv")
```

In [25]:

```
A.shape
```

Out[25]:

```
(150, 6)
```

**Handling categorical values**

In [26]:

```
A["Species"] = A["Species"].map({"Iris-setosa": 0, "Iris-versicolor": 1, "Iris-virginica":
```

**Printing first five and last five records from the dataset**

In [27]:

```
A.head(5)
```

Out[27]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	0
1	2	4.9	3.0	1.4	0.2	0
2	3	4.7	3.2	1.3	0.2	0
3	4	4.6	3.1	1.5	0.2	0
4	5	5.0	3.6	1.4	0.2	0

Type *Markdown* and LaTeX:  $\alpha^2$

In [28]:

```
A.tail(5)
```

Out[28]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
145	146	6.7	3.0	5.2	2.3	2
146	147	6.3	2.5	5.0	1.9	2
147	148	6.5	3.0	5.2	2.0	2
148	149	6.2	3.4	5.4	2.3	2
149	150	5.9	3.0	5.1	1.8	2

## Printing the datatype of each attribute

In [29]:

```
A.dtypes
```

Out[29]:

```
Id                int64
SepalLengthCm     float64
SepalWidthCm      float64
PetalLengthCm     float64
PetalWidthCm      float64
Species           int64
dtype: object
```

## Setting the dependent and independent variables

In [30]:

```
X = A[["SepalLengthCm", "SepalWidthCm", "PetalLengthCm", "PetalWidthCm"]]
```

In [31]:

```
Y = A[["Species"]]
```

## Splitting the dataset into training data and testing data

In [32]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3)
```

## Training the logistic regression model

In [33]:

```
from sklearn.linear_model import LogisticRegression
```

In [34]:

```
reg = LogisticRegression()
```

In [35]:

```
model = reg.fit(X_train, Y_train)
```

C:\Users\RASIKA\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils\validation.py:1111: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

C:\Users\RASIKA\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear\_model\\_logistic.py:444: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

In [36]:

```
Y_predict = model.predict(X_test)
```

In [37]:

```
print(Y_predict)
```

```
[0 0 1 1 1 0 0 2 1 1 2 0 2 0 1 0 2 2 0 1 0 1 0 2 0 2 0 2 2 0 1 0 0 2 2 1 0
 2 2 1 1 2 0 1 1]
```

## Printing the confusion matrix

In [38]:

```
from sklearn.metrics import confusion_matrix
```

In [39]:

```
cm = confusion_matrix(Y_test, Y_predict)
```

In [40]:

```
print(cm)
```

```
[[17  0  0]
 [ 0 14  2]
 [ 0  0 12]]
```

## Printing the classification report

In [41]:

```
from sklearn.metrics import classification_report
```

In [42]:

```
cr = classification_report(Y_test, Y_predict)
```

In [43]:

```
print(cr)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	17
1	1.00	0.88	0.93	16
2	0.86	1.00	0.92	12
accuracy			0.96	45
macro avg	0.95	0.96	0.95	45
weighted avg	0.96	0.96	0.96	45

## Printing the accuracy

In [44]:

```
from sklearn.metrics import accuracy_score  
acc = accuracy_score(Y_test, Y_predict)  
print(acc)
```

0.9555555555555556